

Project_final

April 3, 2021

1 Machine Learning Assignment

```
[30]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
warnings.simplefilter('ignore')
from matplotlib import colors
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split, StratifiedKFold,
    ↳cross_val_score, validation_curve, GridSearchCV
from sklearn.neighbors import KNeighborsClassifier as kNN
from sklearn.linear_model import LogisticRegression as LogReg
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis as QDA
from sklearn.tree import DecisionTreeClassifier as DTreeClass
from sklearn.metrics import accuracy_score, zero_one_loss, roc_auc_score,
    ↳roc_curve, classification_report, f1_score
from scipy import stats
import jupyterthemes
!jt -t grade3 -tfs 7 -nfs 7 -ofs 7 -fs 7
```

```
[31]: names = ['Sequence',
    ↳name', 'mcg', 'gvh', 'alm', 'mit', 'erl', 'pox', 'vac', 'nuc', 'class'];

data_yeast = pd.read_csv('yeast.data', header = None, sep = '\s+', names =
    ↳names);

data_pima = pd.read_csv('diabetes.csv');
```

1.1 Question 1

```
[32]: cols = data_pima.columns[:8]
fig, ax = plt.subplots(2,4, sharex=False, sharey=False, figsize=(20, 8))
```

```

txt="Figure 1. Box plots for the values of the variables according to target_
↪class."
plt.figtext(0.5, 0.01, txt, wrap=True, horizontalalignment='center',
↪fontsize=14)
for i in range(len(cols)):
    sns.boxplot(ax=ax[i//4, i%4], y=cols[i], x='Outcome', data=data_pima)
    ax[i//4, i%4].set(xlabel=cols[i])
plt.show();

```

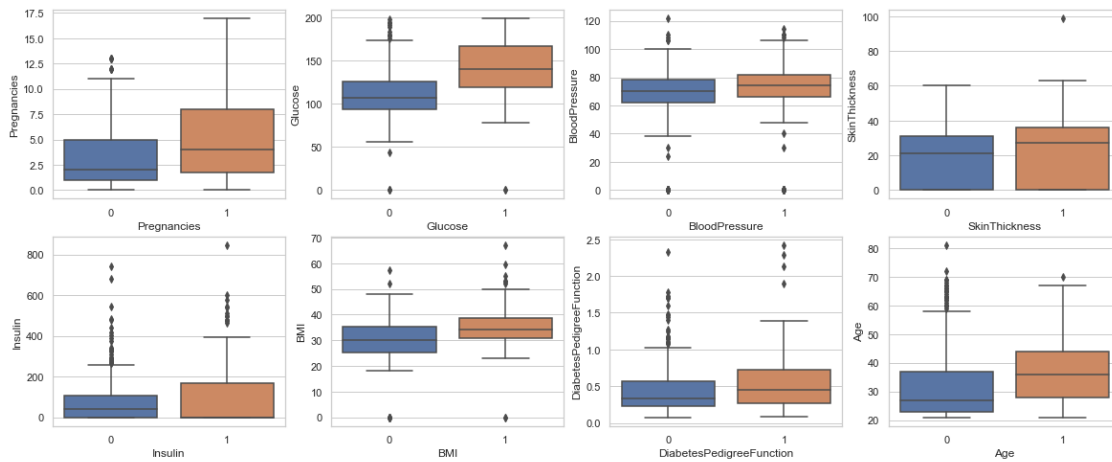
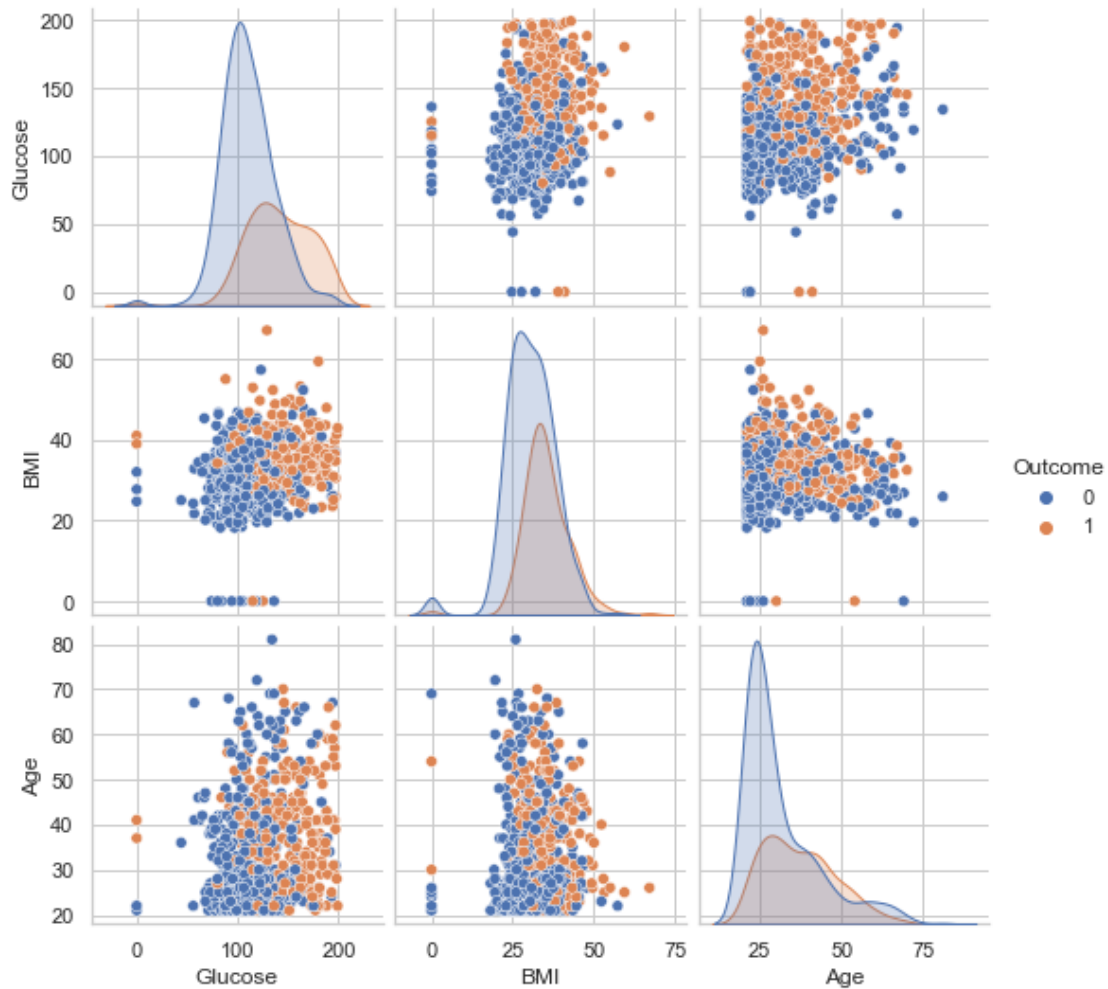


Figure 1. Box plots for the values of the variables according to target class.

```

[33]: data_pima_sel = data_pima.drop(columns = ['BloodPressure', 'SkinThickness',
↪'Pregnancies', 'DiabetesPedigreeFunction', 'Insulin'])
sns.pairplot(data_pima_sel, hue = 'Outcome');

```



Após análise aos gráficos apresentados, podemos observar rapidamente que as variáveis que melhor “separam” a variável alvo são a Glucose, BMI e Age. Contudo, as duas escolhidas (como requisitado pelo exercício) foram a Glucose e BMI. A variável Age foi descartada pois tinha mais outliers.

1.2 Methods comparison (kNN, Logistic Regression, QDA)

```
[34]: data_pima_clean = data_pima_sel[(data_pima['Glucose'] != 0) & (data_pima['BMI'] != 0)]
```

```
data_pima_clean = data_pima_clean.sample(frac=1).reset_index(drop=True)
```

```
[35]: data_pima_sel_final = data_pima_clean.drop(columns = 'Age')
```

```
column_names = ['Glucose', 'BMI']
temp = data_pima_sel_final[column_names].values
temp_scaled = MinMaxScaler().fit_transform(temp)
```

```

data_pima_sel_final[column_names] = temp_scaled

X = data_pima_sel_final.drop(columns = 'Outcome').values

y = data_pima_sel_final.loc[:, 'Outcome'].values
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.20,
    ↪random_state=123, stratify= y)

```

```

[36]: def plot_classifier_boundary(model,x,h = .05): #kindly provided in class
    cmap_light = colors.ListedColormap(['lightsteelblue', 'peachpuff'])
    x_min, x_max = x[:, 0].min()-.2, x[:, 0].max()+.2
    y_min, y_max = x[:, 1].min()-.2, x[:, 1].max()+.2
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
        np.arange(y_min, y_max, h))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)
    plt.contourf(xx, yy, Z, cmap=cmap_light)
    plt.xlim((x_min,x_max))
    plt.ylim((y_min,y_max))
    cmap = colors.ListedColormap(['blue','orange'])

```

```

[37]: def StratKF(model, splits,x_train, y_train, x_test, y_test): #fitted
    ↪model necessary
    SKF = StratifiedKFold(n_splits=splits, shuffle=True, random_state=1)
    score = []
    y_pred_train = model.predict(x_train)
    y_pred_test = model.predict(x_test)
    for train_index, test_index in SKF.split(x_test, y_test):
        x_test1 = x_train[test_index]
        y_test1 = y_train[test_index]
        score.append(f1_score(y_test1, model.predict(x_test1)))
    print(f'F1 score on {splits}-fold test data: ',round(np.mean(score),4),'+/
    ↪-', round(np.std(score),4))
    print('F1 score on training set: ',round(f1_score(y_train,
    ↪y_pred_train),4), '\nF1 score on test set: ',round(f1_score(y_test,
    ↪y_pred_test), 4))

```

1.2.1 Logistic Regression

```

[38]: modelLogReg = LogReg()
    modelLogReg.fit(x_train, y_train)
    StratKF(modelLogReg,5, x_train, y_train, x_test, y_test)

```

```

F1 score on 5-fold test data:  0.524 +/- 0.1142
F1 score on training set:  0.6067
F1 score on test set:  0.5116

```

1.2.2 Quadratic Discriminant Analysis

```
[39]: modelQDA = QDA()
modelQDA.fit(x_train, y_train)
StratKF(modelQDA,5, x_train, y_train, x_test, y_test)
```

F1 score on 5-fold test data: 0.5145 +/- 0.0953


F1 score on training set: 0.6458

F1 score on test set: 0.5169

1.2.3 kNN

```
[40]: fig = plt.figure(figsize=(30,10))
k = [1,10,50,100,200,300]

for i in range(len(k)):
    modelkNN = kNN(n_neighbors=k[i])
    modelkNN.fit(x_train, y_train)
    ax = fig.add_subplot(2, 3, i+1)
    plot_classifier_boundary(modelkNN, X)
    ax.scatter(X[:,0],X[:,1],color=cmap(y))
    ax.set_title(k[i], fontsize = 18)
    ax.set_xlabel('Glucose')
    ax.set_ylabel('BMI');
```



```
[41]: def valid_curve(model, x_train, y_train, n_jobs, scoring, param_range = np.
    ↳ array([*range(1,301)]), param_name="n_neighbors"):
    np.random.seed(1)
    train_scores, test_scores = validation_curve(model, x_train, y_train,
    ↳ param_name=param_name, param_range=param_range, scoring=scoring,
    ↳ n_jobs=n_jobs)
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
```

```

test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

plt.figure(figsize=(20, 8))
plt.title("kNN Validation Curve")
plt.xlabel("k Neighbors")
plt.ylabel("Score")
plt.plot(param_range, train_scores_mean, label="Training score",
↪color="darkorange")
plt.fill_between(param_range, train_scores_mean - train_scores_std,
↪train_scores_mean + train_scores_std, alpha=0.1, color="darkorange")
plt.plot(param_range, test_scores_mean, label="Cross-validation score",
↪color="navy")
plt.fill_between(param_range, test_scores_mean - test_scores_std,
↪test_scores_mean + test_scores_std, alpha=0.1, color="navy")
plt.legend(loc="best")
plt.show();


print("Best K is %d" %param_range[np.where(test_scores_mean ==
↪max(test_scores_mean))][0])
return param_range[np.where(test_scores_mean == max(test_scores_mean))][0]

```

```

[42]: #Determination of k
modelkNN = kNN()
k_best = valid_curve(modelkNN, x_train, y_train, 6, 'accuracy')

```



Best K is 37

```

[43]: modelkNN = kNN(n_neighbors=k_best)
modelkNN.fit(x_train, y_train)
StratKF(modelkNN,5, x_train, y_train, x_test, y_test)

```

F1 score on 5-fold test data: 0.5293 +/- 0.1105
F1 score on training set: 0.6273
F1 score on test set: 0.5227

```
[44]: #ROC AUC
ns_probs = [0 for _ in range(len(y_test))]
# probabilities for the positive outcome
LogReg_probs = modelLogReg.predict_proba(x_test)[: , 1]
QDA_probs = modelQDA.predict_proba(x_test)[: , 1]
kNN_probs = modelkNN.predict_proba(x_test)[: , 1]
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
LogReg_auc = roc_auc_score(y_test, LogReg_probs)
QDA_auc = roc_auc_score(y_test, QDA_probs)
kNN_auc = roc_auc_score(y_test, kNN_probs)
# summarize scores
#print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('LR ROC AUC=%.3f' % (LogReg_auc))
print('QDA ROC AUC=%.3f' % (QDA_auc))
print('kNN ROC AUC=%.3f' % (kNN_auc))
```

LR ROC AUC=0.751
QDA ROC AUC=0.746
kNN ROC AUC=0.759

```
[45]: #ROC curves
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs, pos_label=None)
LogReg_fpr, LogReg_tpr, _ = roc_curve(y_test, LogReg_probs)
QDA_fpr, QDA_tpr, _ = roc_curve(y_test, QDA_probs)
kNN_fpr, kNN_tpr, _ = roc_curve(y_test, kNN_probs)

figs = plt.figure(figsize=(30,8))
# plot the roc curves
axs = figs.add_subplot(1, 5, 1)
axs.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
axs.plot(LogReg_fpr, LogReg_tpr, marker='+', label='Logistic Regression')
axs.plot(QDA_fpr, QDA_tpr, marker='+', label='QDA')
axs.plot(kNN_fpr, kNN_tpr, marker='+', label='kNN')
axs.set_xlabel('False Positive Rate')
axs.set_ylabel('True Positive Rate')
axs.set_title('ROC curves', fontsize = 18)
axs.legend();
##Barplot of AUCs
axs = figs.add_subplot(1, 5, 2)
sns.barplot(y=[LogReg_auc, QDA_auc, kNN_auc], x=['LogReg', 'QDA', 'kNN'], ax=axs)
axs.set_ylim(0,1)
```

```


axs.set_title('ROC AUC Comparison', fontsize = 18)
axs.set_ylabel('ROC AUC');

model_list = [(modelLogReg, 'Logistic Regression Boundary') , (modelQDA, 'QDA_
↳Boundary'), (modelkNN, 'kNN Boundary')]

for i in range(3,6,1):
    axs = figs.add_subplot(1, 5, i)
    plot_classifier_boundary(model_list[i-3][0],X)
    axs.scatter(X[:,0],X[:,1],color=cmap(y))
    axs.set_xlabel('Glucose')
    axs.set_ylabel('BMI');
    axs.set_title(model_list[i-3][1], fontsize = 18)

plt.show();

```



1.3 Question 2

1.3.1 normalization

```

[46]: MinMaxSca = MinMaxScaler()
data_yeast[['mcg', 'gvh', 'alm', 'mit', 'erl', 'pox', 'vac', 'nuc']] =
↳MinMaxSca.fit_transform(data_yeast.loc[:,['mcg', 'gvh', 'alm', 'mit', 'erl',
↳'pox', 'vac', 'nuc']])

```

1.3.2 Separation in train data and test data


```
[47]: X = data_yeast.loc[:, ['mcg', 'gvh', 'alm', 'mit', 'erl', 'pox', 'vac', 'nuc']].
      ↪values
      Y = data_yeast.loc[:, 'class'].values

      X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
      ↪random_state = 1, stratify = Y, shuffle=True)

      SKF = StratifiedKFold(n_splits=5, shuffle=True, random_state=1)
```

```
[48]: def StratKFW(model, splits, X_train, Y_train, X_test, Y_test):
      Y_pred_trainLogReg = model.predict(X_train)
      Y_pred_testLogReg = model.predict(X_test)
      print('Weighted F1 score on training set: ', round(f1_score(Y_train,
      ↪Y_pred_trainLogReg, average='weighted'), 4),
            '\nWeighted F1 score on test set: ', round(f1_score(Y_test,
      ↪Y_pred_testLogReg, average='weighted'), 4))
      f1score = []
      accscore = []
      for train_index, test_index in SKF.split (X_test, Y_test):
          X_test1 = X_test[test_index]
          Y_test1 = Y_test[test_index]
          f1score.append(f1_score(Y_test1, model.predict(X_test1),
      ↪average='weighted'))
          accscore.append(accuracy_score(Y_test1, model.predict(X_test1)))
      print(f'Weighted F1 score on {splits}-fold test data: ', round(np.
      ↪mean(f1score_LogReg), 4), '+/-', round(np.std(f1score_LogReg), 4))
      print('\nClassification report:\n', classification_report(Y_test,
      ↪Y_pred_testLogReg, digits=3))
      return f1score, accscore
```

1.3.3 Logistic regression

```
[49]: modelLR = LogReg()
      modelLR.fit(X_train, Y_train)
      f1score_LogReg, accscore_LogReg = [], []
      f1score_LogReg, accscore_LogReg = StratKFW(modelLR, 5, X_train, Y_train,
      ↪X_test, Y_test)
```

Weighted F1 score on training set: 0.5564

Weighted F1 score on test set: 0.5406

Weighted F1 score on 5-fold test data: nan +/- nan

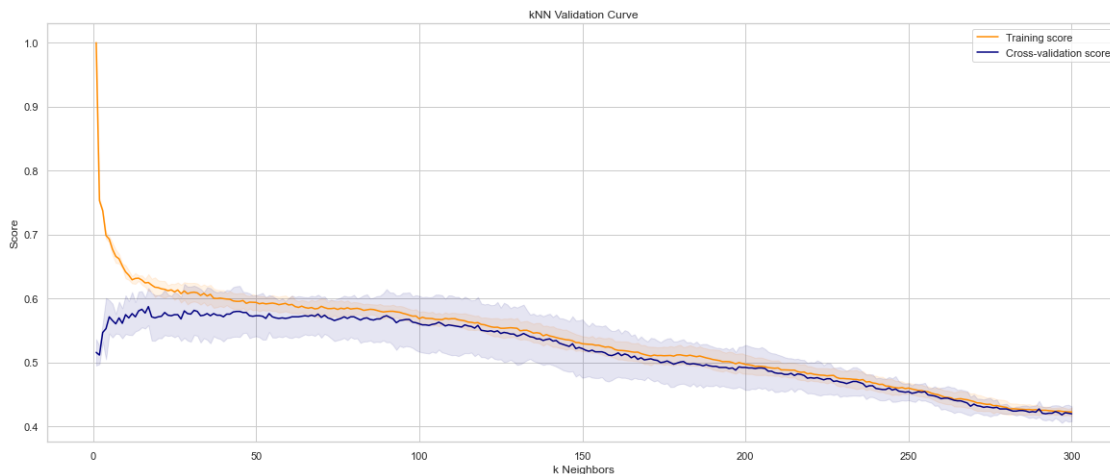
Classification report:

	precision	recall	f1-score	support
CYT	0.504	0.742	0.600	93
ERL	0.000	0.000	0.000	1

EXC	0.000	0.000	0.000	7
ME1	0.455	0.556	0.500	9
ME2	0.000	0.000	0.000	10
ME3	0.667	0.688	0.677	32
MIT	0.617	0.592	0.604	49
NUC	0.631	0.477	0.543	86
POX	0.667	0.500	0.571	4
VAC	0.000	0.000	0.000	6
accuracy			0.566	297
macro avg	0.354	0.355	0.350	297
weighted avg	0.537	0.566	0.541	297

1.3.4 kNN

```
[50]: k_best = valid_curve(modelknn, X_train, Y_train, 6, 'accuracy')
```



Best K is 17

```
[51]: modelknn = kNN(n_neighbors=k_best)
modelknn.fit(X_train, Y_train)
f1score_knn, accscore_knn = [], []
f1score_knn, accscore_knn = StratKFW(modelknn, 5, X_train, Y_train, X_test,
↪Y_test)
```

Weighted F1 score on training set: 0.6207

Weighted F1 score on test set: 0.5658

Weighted F1 score on 5-fold test data: 0.5383 +/- 0.0319

Classification report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

CYT	0.526	0.656	0.584	93
ERL	0.000	0.000	0.000	1
EXC	0.667	0.571	0.615	7
ME1	0.462	0.667	0.545	9
ME2	0.333	0.200	0.250	10
ME3	0.719	0.719	0.719	32
MIT	0.681	0.653	0.667	49
NUC	0.554	0.477	0.513	86
POX	0.667	0.500	0.571	4
VAC	0.000	0.000	0.000	6
accuracy			0.576	297
macro avg	0.461	0.444	0.446	297
weighted avg	0.565	0.576	0.566	297

1.3.5 Decision tree

```
[52]: modeltree = DTreeClass()
modeltree.fit(X_train, Y_train)
f1score_tree, accscore_tree = [], []
f1score_tree, accscore_tree = StratKFW(modeltree, 5, X_train, Y_train, X_test,
→Y_test)
```

Weighted F1 score on training set: 1.0

Weighted F1 score on test set: 0.4973

Weighted F1 score on 5-fold test data: 0.5383 +/- 0.0319

Classification report:

	precision	recall	f1-score	support
CYT	0.515	0.548	0.531	93
ERL	0.000	0.000	0.000	1
EXC	0.300	0.429	0.353	7
ME1	0.727	0.889	0.800	9
ME2	0.111	0.100	0.105	10
ME3	0.688	0.688	0.688	32
MIT	0.556	0.408	0.471	49
NUC	0.489	0.500	0.494	86
POX	0.000	0.000	0.000	4
VAC	0.000	0.000	0.000	6
accuracy			0.498	297
macro avg	0.339	0.356	0.344	297
weighted avg	0.501	0.498	0.497	297

```
[53]: def model_grid_search(model, param_grid, cv, scoring, n_jobs): # "cv - integer,
→to specify the number of folds in a `(Stratified) KFold`, "
    model_grid = GridSearchCV(estimator = model, param_grid = param_grid, cv =
→cv, refit = True, scoring= scoring, n_jobs = n_jobs)
    model_grid.fit(X_train, Y_train)
    print(f"Best estimator: {model_grid.best_estimator_} \n Best score:
→{model_grid.best_score_} \n Best Params: {model_grid.best_params_}")
    return model_grid.best_estimator_
```

```
[54]: max_depth = [None] + [x for x in np.arange(1,20,4)]
min_sample_split= np.arange(2, 10,2)
min_sample_leaf = np.arange(1,5)
ccp_alpha = np.arange(0.01,100, 10)

param_grid_tree = {"criterion": ['gini', "entropy"],\
                    "splitter": ['best',"random"],\
                    "max_depth": max_depth,\
                    "min_samples_split": min_sample_split,\
                    "min_samples_leaf": min_sample_leaf, \
                    "ccp_alpha": ccp_alpha}
```

```
[55]: tree_grid = model_grid_search(modeltree, param_grid_tree, 5, "f1_weighted", 2)
```

```
Best estimator: DecisionTreeClassifier(ccp_alpha=0.01, criterion='entropy',
max_depth=5,
                                min_samples_leaf=4, min_samples_split=4)
Best score: 0.5669393207369835
Best Params: {'ccp_alpha': 0.01, 'criterion': 'entropy', 'max_depth': 5,
'min_samples_leaf': 4, 'min_samples_split': 4, 'splitter': 'best'}
```

```
[56]: modeltree = tree_grid
f1score_tree, accscore_tree = StratKFW(modelLR, 5, X_train, Y_train, X_test,
→Y_test)
```

```
Weighted F1 score on training set: 0.5564
Weighted F1 score on test set: 0.5406
Weighted F1 score on 5-fold test data: 0.5383 +/- 0.0319
```

Classification report:

	precision	recall	f1-score	support
CYT	0.504	0.742	0.600	93
ERL	0.000	0.000	0.000	1
EXC	0.000	0.000	0.000	7
ME1	0.455	0.556	0.500	9
ME2	0.000	0.000	0.000	10
ME3	0.667	0.688	0.677	32

MIT	0.617	0.592	0.604	49
NUC	0.631	0.477	0.543	86
POX	0.667	0.500	0.571	4
VAC	0.000	0.000	0.000	6
accuracy			0.566	297
macro avg	0.354	0.355	0.350	297
weighted avg	0.537	0.566	0.541	297

1.3.6 Achieved results comparison

```
[57]: stats_f1, pvalue_f1 = stats.f_oneway(f1score_LogReg, f1score_knn, f1score_tree)
stats_acc, pvalue_acc = stats.f_oneway(accscore_LogReg, accscore_knn,
↪accscore_tree)
```

```
[58]: fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(15, 5), sharey=True)
sns.set_theme(style="whitegrid")
ax[0].set_title('Models F1 Score', fontsize = 20)
ax[1].set_title('Models Accuracy Score', fontsize = 20)
sns.boxplot(y=f1score_LogReg + f1score_knn + f1score_tree,
↪x=['LogReg']*5+['kNN']*5+['Tree']*5, ax=ax[0])
sns.boxplot(y=accscore_LogReg + accscore_knn + accscore_tree,
↪x=['LogReg']*5+['kNN']*5+['Tree']*5, ax=ax[1]);
```

